



SOUTHERN UNIVERSITY OF SCIENCE AND
TECHNOLOGY

CS212 INTELLIGENT DATA ANALYSIS

Director: Prof. Peter Tiño

Investigating Properties of Houses from a New York Airbnb Dataset

Author:

Edward Yidong FANG

Student Number:

11510493

Email: 11510493@mail.sustc.edu.cn

Shenzhen, China July 6, 2017

Abstract

In these experiments, three basic methods of data analysis are implemented, which is Principle Components Analysis (PCA), Clustering and Self-organizing Map (SOM). By applying these methods, the relations among different attributes are discussed. After the dimension reduction, it was found that the attributes can be predicted by others to some extent, while some cannot. The relations will be presented mostly as different figures.

Contents

1	Introduction	2
2	Data Preprocessing	2
2.1	Data Description	2
2.2	Data Preprocessing	3
2.2.1	Removing of the useless data and the Selection of the records	3
2.2.2	One Hot Encoding	4
2.2.3	Data Standardization	4
3	Questions to Ask and Labeling	6
3.1	Questions to Ask	6
3.2	Labeling Schemes	7
3.2.1	Label of Prices	7
3.2.2	Label of Overall Satisfaction	7
4	Data Analysis	8
4.1	Coordinate Projections	8
4.2	Principal Component Analysis	8
4.2.1	Prices of the rooms	8
4.2.2	Overall satisfaction	10
4.3	Clustering	10
4.4	Self-organization map	10
4.5	Differences between two boroughs	11
5	Conclusion	18

1 Introduction

Airbnb is an online marketplace and hospitality service, enabling people to lease or rent short-term lodging including vacation rentals, apartment rentals, homestays, hostel beds, or hotel rooms. The dataset analysed in there experiments is derived from Tom Slee’s blog[1] and it is crawled from the website of Airbnb. And only a small part of data for New York, crawled on 05/06/2017, are analysed.

2 Data Preprocessing

2.1 Data Description

Here is the meaning for each column in the collected CSV file:

- `room_id`: A unique number identifying an Airbnb listing. The listing has a URL on the Airbnb web site of http://airbnb.com/rooms/room_id
- `survey_id`: A unique number identifying the behaviour of survey.
- `host_id`: A unique number identifying an Airbnb host. The host’s page has a URL on the Airbnb web site of http://airbnb.com/users/show/host_id
- `room_type`: One of “Entire home/apt”, “Private room”, or “Shared room”
- `country`: the nation the room located in; acutually no data
- `city`: the city the room located in
- `borough`: A subregion of the city or search area for which the survey is carried out. The borough is taken from a shapefile of the city that is obtained independently of the Airbnb web site. For some cities, there is no borough information; for others the borough may be a number. If you have better shapefiles for a city of interest, please send them to me.
- `neighborhood`: As with `borough`: a subregion of the city or search area for which the survey is carried out. For cities that have both, a neighbourhood is smaller than a borough. For some cities there is no neighbourhood information.
- `reviews`: The number of reviews that a listing has received. Airbnb has said that 70% of visits end up with a review, so the number of reviews can be used to estimate the number of visits. Note that such an estimate will not be reliable for an individual listing (especially as reviews occasionally vanish from the site), but over a city as a whole it should be a useful metric of traffic.

- `overall_satisfaction`: The average rating (out of five) that the listing has received from those visitors who left a review.
- `accommodates`: The number of guests a listing can accommodate.
- `bedrooms`: The number of bedrooms a listing offers.
- `bathrooms`: The number of bathrooms a listing offers, actually not available.
- `price`: The price (in \$US) for a night stay.
- `minstay`: The minimum stay for a visit, as posted by the host.
- `name`: The name of the room.
- `property_type`: “Apartment”, “Loft”, “Villa”, “House”, etc.
- `latitude` and `longitude`: The latitude and longitude of the listing as posted on the Airbnb site: this may be off by a few hundred metres.
- `last_modified`: the date and time that the values were read from the Airbnb web site.
- `location`: Unknown, certain number related to the location of the room.

The first line of the CSV file holds the column headings.

2.2 Data Preprocessing

Note that in this report only the data preprocessing process of the first problem is shown, but the other data preprocessing processes are all very similar.

2.2.1 Removing of the useless data and the Selection of the records

As the amount of data is really large (up to 40,730 rows), we just remove the data records without available *reviews* or *overall_satisfaction*. Since the prices higher than \$200 per night are assumed as not reasonable, the records with those extremely high prices are removed. Also, as the columns *room_id*, *survey_id*, *host_id*, *country*, *city*, *last_modified* and *location* are meaningless, we removed them from the data. In addition, it is found that the values of column *bathrooms* and *minstay* are unavailable.

Finally, for the simplification of the problem, we just retrieve the records whose boroughs are “Manhattan” or “Brooklyn”. And the experimental data are sampled from the original data under the sample fraction of 0.1.

	room_type	borough	accommodates	reviews	overall_satisfaction	bedrooms	price	longitude	latitude	property_type
0	Entire home/apt	Brooklyn	2	95	5.0	0.0	125.0	-73.943276	40.721256	Apartment
1	Entire home/apt	Brooklyn	3	3	4.5	1.0	165.0	-73.952168	40.723975	House
2	Entire home/apt	Manhattan	5	25	4.0	2.0	220.0	-73.962862	40.758275	Apartment
3	Entire home/apt	Manhattan	2	2	0.0	1.0	180.0	-74.003905	40.733196	Apartment
4	Entire home/apt	Brooklyn	4	3	5.0	1.0	132.0	-73.957279	40.733538	Apartment

Table 1: Data After preprocessing

The data at the end looks like the Table 1.

2.2.2 One Hot Encoding

To make use of the data field *room_type*, *borough* and *property_type*, both of the information are encoded using the method called One-Hot encoding to transform the data. By doing so, it is ensured that the information of all possible aspects of the room are included into the data matrix.

For example, if there two room, whose types are “Private room” and “Entire home/apt”, respectively, then the encoded result looks like Table 2. Three different columns are added to enumerate all possible room types. And the values are all zeros except for the room type the room belongs to.

	accommodates	...	latitude	room_type_Entire home/apt	room_type_Private room	room_type_Shared room	...
0	2	...	40.691398	0	1	0	...
1	4	...	40.811016	1	0	0	...

Table 2: Example Data After One-hot Encoding

2.2.3 Data Standardization

To see the significance of data standardization, two box plot was made, one for the raw data and another for the standardized data. Note that the one-hot encoded columns are standardized but are not included in the figures because the scale of the box plot will be influenced by those attributes and make the difference hard to distinguish.

As is obvious from the Figure 1, the values of attribute *reviews* are not in the same scale with other attributes. Neither does the longitude and latitude. If we use these values in the matrix we will calculate later, the importance of attribute *reviews*, *latitude* and *longitude* will be extremely larger than other attributes. Thus, we need to apply the standardization to the raw data.




Figure 1: Data before standardization

Standardization can be done in many ways, the formula

$$NewValue = (OldValue - Mean) / StandardDeviation$$

is enough for this dataset. The result is shown in the Figure 2.




Figure 2: Data after standardization

3 Questions to Ask and Labeling

3.1 Questions to Ask

When people go out and want to find a room to stay in a city, it can be said that the most concerning aspects of a room is its condition, position and price. In the dataset analysed, the price is already there and the overall satisfaction can represent a combination of all the aspects. Thus, two questions are going to be explored in these experiment.

- Can the ranges of price be roughly predicted using other attributes of the data?
- Can the overall satisfactions be roughly predicted using other attributes of data?
- What is the most crucial factors that influence the price and overall satisfaction among the attributes we have?
- How much the difference will be for the data in two different borough?

3.2 Labeling Schemes

3.2.1 Label of Prices




Figure 3: Distribution of prices that is lower than \$750

The price level can be classified into three levels according to the distribution of the prices. Fig.3 shows the distribution and according to it, the data are labelled into three approximately-equal-size classes: “low price”, “middle price” and “high price”.

Note that for the convenience of visualization, we just plot the distribution of prices that is lower than \$750, and only a few records’ prices are higher than \$750.

After some trying, it is finally found that these three intervals are best for the data to be evenly divided and labelled: $[0, 85)$, $[85, 149)$ and $[149, \infty)$.

3.2.2 Label of Overall Satisfaction

The label of overall satisfaction is set to be “low” if *overall_satisfaction* is equal to or less than 4, “median” if it is equal to 4.5, “high” if equal to 5.

4 Data Analysis

After the preprocessing of the data, the Principle Component Analysis (PCA) algorithm was used to generate the eigenvalues, eigenvectors, and coefficients to the eigenvalues, which actually are the variances that contained in each directions that eigenvectors expend.

4.1 Coordinate Projections

First some simple projections of just two attributes of the data are used.

As most of our attributes are discrete, some of the important can be distinguished as a form of points' labels.

From most of these figures such as Fig.4 and Fig.5 it can be found that the average room price in Manhattan is generally higher than that in Brooklyn. This is perhaps due to the fact that Manhattan is the center of the city. Also, Fig.4 indicates that the larger the accomodates number is, the higher the price is. The trend of price over number of bedrooms in Fig.5 is similar, which is intuitively true.

However, Fig. 6, 8, 9, 10 and 11 also show some relation between price and location, overall satisfaction, room type and reviews number. So which of these attributes can influence the price mostly, i.e. the importance of these attributes. Then it can be figured out by the technique of PCA.

4.2 Principal Component Analysis

4.2.1 Prices of the rooms

It can be noticed from Fig.15 that the distribution of variance of the principle components is not concentrated to the first several principal components very well. Thus, if the data is projected on to the first two or three principle components, it can be predicted that the structure of results will not very similar to that of the original data.

However, the PCA technique indeed reduce some dimensions, as it can be concluded from the Fig.15 that the first 15 eigenvectors can preserve the 90% of the variance of the dataset.

Figure 13 projects the 23 dimensional data onto a sub-space spanned by the first and second eigenvector of the covariance matrix of X




Figure 4: A 2D-projection on to the prices and accomdates with borough label
 labelfig:pair-over-price2

The more detailed information of the first three eigenvectors are shown in the Table 3. In the first eigenvector the importance of borough is very important, this is why it can be discovered that in the following PCA projection figures there are 2 similar part of cluster (such as in Fig. 13 and 14). Also, the latitude and longitude are actually related to the borough it located in, thus they increase the distinction or distance in the projection.

The price differences are mainly interpreted by the second principal eigenvector. The most significant components in the 2nd vector is the room type. Since “Entire” type and “Private” type are negative and positive respectively, we can conclude from the PCA projection result that the prices of Entire home/apt rooms are higher much higher than those of Private room. Also, the more accommodates and bedrooms, the higher price the room will have. The importance of accommodate numbers is greater than that of room numbers. And the overall stratification does influence a little on the price, which reveal the “effect of market”.

For the 3rd most significant eigenvector, the point is that the room with larger value has lower price. So this vector illustrates the property type influence on the price. The price of property type from lower to higher is apartment, house, loft and townhouse.




Figure 5: A 2D-projection on to the prices and bedrooms with borough label

1st Principle Component Eigenvalue 3.15901473		2nd Principle Component Eigenvalue 2.53613192		3rd Principle Component Eigenvalue 1.73809388	
borough_Manhattan	-0.51484551	room_type_Entire home/apt	-0.52678505	property_type_Apartment	0.60738014
borough_Brooklyn	0.51484551	room_type_Private room	0.51272818	property_type_House	-0.36991721
latitude	-0.41706656	accommodates	-0.50934697	property_type_Loft	-0.28433712
property_type_Apartment	-0.29132238	bedrooms	-0.34241624	borough_Manhattan	-0.26485932
longitude	0.26209607	overall_satisfaction	-0.14247044	borough_Brooklyn	0.26485932
property_type_House	0.25374459	latitude	0.11644046	latitude	-0.22556836
room_type_Private room	0.15672518	reviews	-0.0980652	property_type_Townhouse	-0.20936499
room_type_Entire home/ap	-0.15512859	property_type_Apartment	0.09209157	room_type_Entirehome/apt	0.1495919
property_type_Loft	0.10969607	longitude	0.08412272	property_type_Other	-0.14047292

Table 3: Details of top three eigenvectors

4.2.2 Overall satisfaction

Then the same technique is applied to the information without overall satisfaction data.

4.3 Clustering

4.4 Self-organization map

Each cell is the normalised sum of the distances between a neuron and its neighbours.




Figure 6: A 2D-projection on to the latitude and longitude with price label

4.5 Differences between two boroughs




Figure 7: A 2D-projection on to the latitude and longitude with satisfiction label




Figure 8: A 2D-projection on to the satisfiction and price with borough label




Figure 9: A 2D-projection on to the satisfaction and price with room type label




Figure 10: A 2D-projection on to the number of reviews and price with borough label




Figure 11: A 2D-projection on to the number of reviews and price with room type label




Figure 12: A scree plot showing the percentage variance of each principle component, and a running cumulative total




Figure 13: A 2D projection of the PCA results




Figure 14: A 3D projection of the PCA results

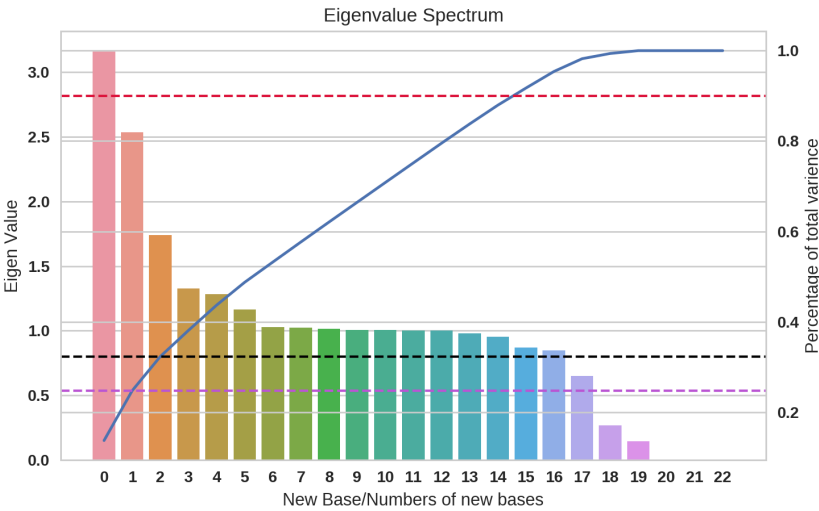


Figure 15: A scree plot showing the percentage variance of each principle component, and a running cumulative total




Figure 16: A 2D projection of the PCA results




Figure 17: A 3D projection of the PCA results

5 Conclusion

References

- [1] T. Slee. Airbnb data collection: Get the data, 2017. URL <http://tomslee.net/airbnb-data-collection-get-the-data>.